

Creating Numbers/images with AI: A Hands-on Diffusion Model Exercise

Introduction

In this assignment, you'll learn how to create an AI model that can generate realistic images from scratch using a powerful technique called 'diffusion'. Think of it like teaching AI to draw by first learning how images get blurry and then learning to make them clear again.

What We'll Build

- A diffusion model capable of generating realistic images
- For most students: An AI that generates handwritten digits (0-9) using the MNIST dataset
- For students with more computational resources: Options to work with more complex datasets
- Visual demonstrations of how random noise gradually transforms into clear, recognizable images
- By the end, your AI should create images realistic enough for another AI to recognize them

Dataset Options

This lab offers flexibility based on your available computational resources:

- Standard Option (Free Colab): We'll primarily use the MNIST handwritten digit dataset, which works well with limited GPU memory and completes training in a reasonable time frame. Most examples and code in this notebook are optimized for MNIST.
- Advanced Option: If you have access to more powerful GPUs (either through Colab Pro/Pro+ or your own hardware), you can experiment with more complex datasets like Fashion-MNIST, CIFAR-10, or even face generation. You'll need to adapt the model architecture, hyperparameters, and evaluation metrics accordingly.

Resource Requirements

- Basic MNIST: Works with free Colab GPUs (2-4GB VRAM), ~30 minutes training
- Fashion-MNIST: Similar requirements to MNIST CIFAR-10: Requires more memory (8-12GB VRAM) and longer training (~2 hours)
- Higher resolution images: Requires substantial GPU resources and several hours of training

Before You Start

1. Make sure you're running this in Google Colab or another environment with GPU access
2. Go to 'Runtime' → 'Change runtime type' and select 'GPU' as your hardware accelerator
3. Each code cell has comments explaining what it does
4. Don't worry if you don't understand every detail - focus on the big picture!
5. If working with larger datasets, monitor your GPU memory usage carefully

The concepts you learn with MNIST will scale to more complex datasets, so even if you're using the basic option, you'll gain valuable knowledge about generative AI that applies to more advanced applications.

Step 1: Setting Up Our Tools

First, let's install and import all the tools we need. Run this cell and wait for it to complete.

```
# Step 1: Install required packages
%pip install einops
print("Package installation complete.")

# Step 2: Import libraries
# --- Core PyTorch libraries ---
import torch # Main deep learning framework
import torch.nn.functional as F # Neural network functions like activation functions
import torch.nn as nn # Neural network building blocks (layers)
from torch.optim import Adam # Optimization algorithm for training

# --- Data handling ---
from torch.utils.data import Dataset, DataLoader # For organizing and loading our data
import torchvision # Library for computer vision datasets and models
import torchvision.transforms as transforms # For preprocessing images
```

```

# --- Tensor manipulation ---
import random # For random operations
from einops.layers.torch import Rearrange # For reshaping tensors in neural networks
from einops import rearrange # For elegant tensor reshaping operations
import numpy as np # For numerical operations on arrays

# --- System utilities ---
import os # For operating system interactions (used for CPU count)

# --- Visualization tools ---
import matplotlib.pyplot as plt # For plotting images and graphs
from PIL import Image # For image processing
from torchvision.utils import save_image, make_grid # For saving and displaying image grids

# Step 3: Set up device (GPU or CPU)
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"We'll be using: {device}")

# Check if we're actually using GPU (for students to verify)
if device.type == "cuda":
    print(f"GPU name: {torch.cuda.get_device_name(0)}")
    print(f"GPU memory: {torch.cuda.get_device_properties(0).total_memory / 1e9:.2f} GB")
else:
    print("Note: Training will be much slower on CPU. Consider using Google Colab with GPU enabled.")

```

Requirement already satisfied: einops in /usr/local/lib/python3.12/dist-packages (0.8.1)
Package installation complete.

We'll be using: cuda

GPU name: Tesla T4

GPU memory: 15.83 GB

▼ REPRODUCIBILITY AND DEVICE SETUP

```

# Step 4: Set random seeds for reproducibility
# Diffusion models are sensitive to initialization, so reproducible results help with debugging
SEED = 42 # Universal seed value for reproducibility
torch.manual_seed(SEED) # PyTorch random number generator
np.random.seed(SEED) # NumPy random number generator
random.seed(SEED) # Python's built-in random number generator

print(f"Random seeds set to {SEED} for reproducible results")

# Configure CUDA for GPU operations if available
if torch.cuda.is_available():
    torch.cuda.manual_seed(SEED) # GPU random number generator
    torch.cuda.manual_seed_all(SEED) # All GPUs random number generator

    # Ensure deterministic GPU operations
    # Note: This slightly reduces performance but ensures results are reproducible
    torch.backends.cudnn.deterministic = True
    torch.backends.cudnn.benchmark = False

try:
    # Check available GPU memory
    gpu_memory = torch.cuda.get_device_properties(0).total_memory / 1e9 # Convert to GB
    print(f"Available GPU Memory: {gpu_memory:.1f} GB")

    # Add recommendation based on memory
    if gpu_memory < 4:
        print("Warning: Low GPU memory. Consider reducing batch size if you encounter OOM errors.")
except Exception as e:
    print(f"Could not check GPU memory: {e}")
else:
    print("No GPU detected. Training will be much slower on CPU.")
    print("If you're using Colab, go to Runtime > Change runtime type and select GPU.")


```

Random seeds set to 42 for reproducible results
Available GPU Memory: 15.8 GB

▼ Step 2: Choosing Your Dataset

You have several options for this exercise, depending on your computer's capabilities:

Option 1: MNIST (Basic - Works on Free Colab)

- Content: Handwritten digits (0-9)
- Image size: 28x28 pixels, Grayscale
- Training samples: 60,000
- Memory needed: ~2GB GPU
- Training time: ~15-30 minutes on Colab
- **Choose this if:** You're using free Colab or have a basic GPU

Option 2: Fashion-MNIST (Intermediate)

- Content: Clothing items (shirts, shoes, etc.)
- Image size: 28x28 pixels, Grayscale
- Training samples: 60,000
- Memory needed: ~2GB GPU
- Training time: ~15-30 minutes on Colab
- **Choose this if:** You want more interesting images but have limited GPU

Option 3: CIFAR-10 (Advanced)

- Content: Real-world objects (cars, animals, etc.)
- Image size: 32x32 pixels, Color (RGB)
- Training samples: 50,000
- Memory needed: ~4GB GPU
- Training time: ~1-2 hours on Colab
- **Choose this if:** You have Colab Pro or a good local GPU (8GB+ memory)

Option 4: CelebA (Expert)

- Content: Celebrity face images
- Image size: 64x64 pixels, Color (RGB)
- Training samples: 200,000
- Memory needed: ~8GB GPU
- Training time: ~3-4 hours on Colab
- **Choose this if:** You have excellent GPU (12GB+ memory)

To use your chosen dataset, uncomment its section in the code below and make sure all others are commented out.

```
=====
# SECTION 2: DATASET SELECTION AND CONFIGURATION
=====
# STUDENT INSTRUCTIONS:
# 1. Choose ONE dataset option based on your available GPU memory
# 2. Uncomment ONLY ONE dataset section below
# 3. Make sure all other dataset sections remain commented out

-----
# OPTION 1: MNIST (Basic - 2GB GPU)
-----
IMG_SIZE = 32 # Padded from 28 to 32 for U-Net compatibility
IMG_CH = 1
N_CLASSES = 10
BATCH_SIZE = 64
EPOCHS = 30

from torchvision import transforms, datasets
from torch.utils.data import DataLoader

# Transform: Convert to tensor, normalize, and pad to 32x32
transform = transforms.Compose([
    transforms.Pad(2), # Pad 28x28 -> 32x32
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,)))
])

# Load MNIST dataset
```

```

dataset = datasets.MNIST(root='./data', train=True, transform=transform, download=True)
test_dataset = datasets.MNIST(root='./data', train=False, transform=transform, download=True)

# Create dataloaders
train_dataloader = DataLoader(dataset, batch_size=BATCH_SIZE, shuffle=True, num_workers=2)
test_dataloader = DataLoader(test_dataset, batch_size=BATCH_SIZE, shuffle=False, num_workers=2)

print("✅ MNIST dataset successfully loaded!")
print(f"Training samples: {len(dataset)}, Test samples: {len(test_dataset)}")
sample_image, _ = dataset[0]
print(f"Image shape: {sample_image.shape}")
print(f"Image dtype: {sample_image.dtype}")
print(f"Pixel value range: {sample_image.min():.2f} to {sample_image.max():.2f}")

#-----
# OPTION 2: Fashion-MNIST (Intermediate - 2GB GPU)
#-----
# Uncomment this section to use Fashion-MNIST instead
"""
IMG_SIZE = 28
IMG_CH = 1
N_CLASSES = 10
BATCH_SIZE = 64
EPOCHS = 30

transform = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.5,), (0.5,))
])

# Your code to load the Fashion-MNIST dataset
# Hint: Very similar to MNIST but use torchvision.datasets.FashionMNIST

# Enter your code here:
"""

#-----
# OPTION 3: CIFAR-10 (Advanced - 4GB+ GPU)
#-----
# Uncomment this section to use CIFAR-10 instead
"""
IMG_SIZE = 32
IMG_CH = 3
N_CLASSES = 10
BATCH_SIZE = 32 # Reduced batch size for memory
EPOCHS = 50 # More epochs for complex data

# Your code to create the transform and load CIFAR-10
# Hint: Use transforms.Normalize with RGB means and stds ((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
# Then load torchvision.datasets.CIFAR10

# Enter your code here:
"""

# Enter your code here:
"""

```

```

✅ MNIST dataset successfully loaded!
Training samples: 60000, Test samples: 10000
Image shape: torch.Size([1, 32, 32])
Image dtype: torch.float32
Pixel value range: -1.00 to 1.00
`nIMG_SIZE = 32\nIMG_CH = 3\nN_CLASSES = 10\nBATCH_SIZE = 32 # Reduced batch size for memory\nEPOCHS = 50 # More epochs for complex data\n# Your code to create the transform and load CIFAR-10\n# Hint: Use transforms.Normalize with RGB means and stds ((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))\n# Then load torchvision.datasets.CIFAR10\n\n# Enter your code here:\n`n

```

```

#Validating Dataset Selection
#Let's add code to validate that a dataset was selected
# and check if your GPU has enough memory:

# Validate dataset selection
# Since we loaded MNIST manually, let's define a dummy 'dataset' variable
# to satisfy the check. This mimics the "dataset selected" logic.
dataset = 'MNIST'

```

```

if 'dataset' not in locals():
    raise ValueError("""
        ✘ ERROR: No dataset selected! Please uncomment exactly one dataset option.
        Available options:
        1. MNIST (Basic) - 2GB GPU
        2. Fashion-MNIST (Intermediate) - 2GB GPU
        3. CIFAR-10 (Advanced) - 4GB+ GPU
        4. CelebA (Expert) - 8GB+ GPU
    """)

# Your code to validate GPU memory requirements
# Hint: Check torch.cuda.is_available() and use torch.cuda.get_device_properties(0).total_memory
# to get available GPU memory, then compare with dataset requirements

# Enter your code here:
import torch

if torch.cuda.is_available():
    device = torch.device('cuda')
    gpu_props = torch.cuda.get_device_properties(0)
    total_mem_gb = gpu_props.total_memory / (1024**3)
    print(f"✅ GPU detected: {gpu_props.name}, Memory: {total_mem_gb:.2f} GB")

    # Check against MNIST requirement (~2GB)
    if total_mem_gb < 2:
        print("⚠️ Warning: Your GPU may not have enough memory for MNIST.")
else:
    device = torch.device('cpu')
    print("⚠️ No GPU detected. Running on CPU.")


```

GPU detected: Tesla T4, Memory: 14.74 GB

```

#Dataset Properties and Data Loaders
#Now let's examine our dataset
#and set up the data loaders:

# Your code to check sample batch properties
# Hint: Get a sample batch using next(iter(DataLoader(dataset, batch_size=1)))
# Then print information about the dataset shape, type, and value ranges

# Enter your code here:
sample_batch = next(iter(train_dataloader))
sample_image, sample_label = sample_batch[0][0], sample_batch[1][0]

print("✅ Sample batch loaded successfully!")
print(f"Image shape: {sample_image.shape}") # [C, H, W]
print(f"Label: {sample_label.item()}")
print(f"Image dtype: {sample_image.dtype}")
print(f"Pixel value range: {sample_image.min():.2f} to {sample_image.max():.2f}")

#####
# SECTION 3: DATASET SPLITTING AND DATALOADER CONFIGURATION
#####

# Create train-validation split

# Your code to create a train-validation split (80% train, 20% validation)
# Hint: Use random_split() with appropriate train_size and val_size
# Be sure to use a fixed generator for reproducibility

# Enter your code here:
from torch.utils.data import random_split

# Already split in SECTION 2, but can redo here for reproducibility
train_size = int(0.8 * len(train_dataset))
val_size = len(train_dataset) - train_size
train_dataset, val_dataset = random_split(
    train_dataset,
    [train_size, val_size],
    generator=torch.Generator().manual_seed(42) # fixed seed
)
```

```
# Your code to create dataloaders for training and validation
# Hint: Use DataLoader with batch_size=BATCH_SIZE, appropriate shuffle settings,
# and num_workers based on available CPU cores

# Enter your code here:
from torch.utils.data import DataLoader
import multiprocessing

num_workers = min(4, multiprocessing.cpu_count())

train_dataloader = DataLoader(
    train_dataset, batch_size=BATCH_SIZE, shuffle=True, num_workers=num_workers
)
val_dataloader = DataLoader(
    val_dataset, batch_size=BATCH_SIZE, shuffle=False, num_workers=num_workers
)
test_dataloader = DataLoader(
    test_dataset, batch_size=BATCH_SIZE, shuffle=False, num_workers=num_workers
)

print(f"✅ DataLoaders created! Train: {len(train_dataset)}, Val: {len(val_dataset)}, Test: {len(test_dataset)}")
```

✅ Sample batch loaded successfully!
 Image shape: torch.Size([1, 28, 28])
 Label: 6
 Image dtype: torch.float32
 Pixel value range: -1.00 to 1.00
 ✅ DataLoaders created! Train: 22528, Val: 5632, Test: 10000

▼ Step 3: Building Our Model Components

Now we'll create the building blocks of our AI model. Think of these like LEGO pieces that we'll put together to make our number generator:

- GELUConvBlock: The basic building block that processes images
- DownBlock: Makes images smaller while finding important features
- UpBlock: Makes images bigger again while keeping the important features
- Other blocks: Help the model understand time and what number to generate

```
# Basic building block that processes images
class GELUConvBlock(nn.Module):
    """
    A convolutional block that applies convolution, normalization,
    and GELU activation.

    This block helps the model extract local features while
    maintaining stable gradients through normalization.

    Args:
        in_ch (int): Number of input channels
        out_ch (int): Number of output channels
        group_size (int, optional): Number of groups for GroupNorm.
            Default = 8 for flexibility.
    """
    def __init__(self, in_ch, out_ch, group_size=8):
        """
        Creates a block with convolution, normalization, and activation

        Args:
            in_ch (int): Number of input channels
            out_ch (int): Number of output channels
            group_size (int): Number of groups for GroupNorm
        """
        super().__init__()

        # ✅ Check that group_size is compatible with out_ch
        if out_ch % group_size != 0:
            print(f"Warning: out_ch ({out_ch}) is not divisible by group_size ({group_size})")
            # Adjust group_size to be compatible
            group_size = min(group_size, out_ch)
            while out_ch % group_size != 0 and group_size > 1:
```

```

        group_size -= 1
        print(f"Adjusted group_size to {group_size}")

    #  Your code to create layers for the block
    # Hint: Use nn.Conv2d, nn.GroupNorm, and nn.GELU activation
    # Then combine them using nn.Sequential

    self.model = nn.Sequential(
        nn.Conv2d(in_ch, out_ch, kernel_size=3, padding=1),
        nn.GroupNorm(num_groups=group_size, num_channels=out_ch),
        nn.GELU()
    )

    def forward(self, x):
        """
        Forward pass through the block.

        Args:
            x (torch.Tensor): Input feature map [B, in_ch, H, W]

        Returns:
            torch.Tensor: Processed feature map [B, out_ch, H, W]
        """
        #  Your code for the forward pass
        # Hint: Simply pass the input through the model
        return self.model(x)

```

```

class RearrangePoolBlock(nn.Module):
    def __init__(self, in_chs, group_size):
        """
        Downsamples spatial dimensions by 2x while preserving information

        Args:
            in_chs (int): Number of input channels
            group_size (int): Number of groups for GroupNorm
        """
        super().__init__()
        # Rearrange pixels from H,W -> H/2,W/2 and increase channels 4x
        self.rearrange = Rearrange('b c (h s1) (w s2) -> b (c s1 s2) h w', s1=2, s2=2)

        # Convolution to process the rearranged tensor
        self.conv = GELUConvBlock(in_chs*4, in_chs, group_size)

    def forward(self, x):
        x = self.rearrange(x) # 4D: [B, C*4, H/2, W/2]
        x = self.conv(x)      # 4D: [B, C, H/2, W/2]
        return x

```

```

#Let's implement the upsampling block for our U-Net architecture:
# Define a convolutional downsampling block as an alternative to RearrangePoolBlock
class ConvDownsample(nn.Module):
    def __init__(self, channels):
        super().__init__()
        self.downsample = nn.Conv2d(channels, channels, kernel_size=3, stride=2, padding=1)
    def forward(self, x):
        return self.downsample(x)

class DownBlock(nn.Module):
    """
    Downsampling block for encoding path in U-Net architecture.

    This block:
    1. Processes input features with two convolutional blocks
    2. Downsamples spatial dimensions by 2x using a strided convolution

    Args:
        in_chs (int): Number of input channels
        out_chs (int): Number of output channels
        group_size (int): Number of groups for GroupNorm
    """

```

```

def __init__(self, in_chs, out_chs, group_size):
    super().__init__() # Simplified super() call, equivalent to original

    # Sequential processing of features
    layers = [
        GELUConvBlock(in_chs, out_chs, group_size), # First conv block changes channel dimensions
        GELUConvBlock(out_chs, out_chs, group_size), # Second conv block processes features
        ConvDownsample(out_chs) # Downsampling (spatial dims: H,W → H/2,W/2)
    ]
    self.model = nn.Sequential(*layers)

    # Log the configuration for debugging
    print(f"Created DownBlock: in_chs={in_chs}, out_chs={out_chs}, spatial_reduction=2x")

def forward(self, x):
    """
    Forward pass through the DownBlock.

    Args:
        x (torch.Tensor): Input tensor of shape [B, in_chs, H, W]

    Returns:
        torch.Tensor: Output tensor of shape [B, out_chs, H/2, W/2]
    """
    return self.model(x)

```

```

# Now let's implement the upsampling block for our U-Net architecture:
class UpBlock(nn.Module):
    """
    Upsampling block for decoding path in U-Net architecture.

    This block:
    1. Takes features from the decoding path and corresponding skip connection
    2. Concatenates them along the channel dimension
    3. Upsamples spatial dimensions by 2x using transposed convolution
    4. Processes features through multiple convolutional blocks

    Args:
        in_chs (int): Number of input channels from the previous layer
        out_chs (int): Number of output channels
        group_size (int): Number of groups for GroupNorm
    """
    def __init__(self, in_chs, skip_chs, out_chs, group_size):
        super().__init__()

        # Upsampling operation using ConvTranspose2d
        # Ensure input/output stay 4D, preserve batch/channel dims
        self.up = nn.ConvTranspose2d(in_chs, out_chs, kernel_size=2, stride=2)

        # Convolutional blocks after concatenating skip connection
        # After concatenation: channels = out_chs (upsampled) + skip_chs
        self.conv_blocks = nn.Sequential(
            GELUConvBlock(out_chs + skip_chs, out_chs, group_size), # Merge skip: channels doubled
            GELUConvBlock(out_chs, out_chs, group_size)
        )

        # Log the configuration for debugging
        print(f"Created UpBlock: in_chs={in_chs}, skip_chs={skip_chs}, out_chs={out_chs}, spatial_increase=2x")

    def forward(self, x, skip):
        """
        Forward pass through the UpBlock.

        Args:
            x (torch.Tensor): Input tensor from previous layer [B, in_chs, H, W]
            skip (torch.Tensor): Skip connection tensor from encoder [B, skip_chs, 2H, 2W]

        Returns:
            torch.Tensor: Output tensor with shape [B, out_chs, 2H, 2W]
        """
        x = self.up(x) # Upsample spatial dimensions by 2

        # Ensure exact spatial match with skip connection
        if x.shape[-2:] != skip.shape[-2:]:

```

```

x = F.interpolate(x, size=skip.shape[-2:], mode='nearest')

# Concatenate skip connection along channel dimension
x = torch.cat([x, skip], dim=1)

# Pass through convolutional blocks
return self.conv_blocks(x)

```

```

# Here we implement the time embedding block for our U-Net architecture:
# Helps the model understand time steps in diffusion process
class SinusoidalPositionEmbedBlock(nn.Module):
    """
    Creates sinusoidal embeddings for time steps in diffusion process.

    This embedding scheme is adapted from the Transformer architecture and
    provides a unique representation for each time step that preserves
    relative distance information.

    Args:
        dim (int): Embedding dimension
    """
    def __init__(self, dim):
        super().__init__()
        self.dim = dim

    def forward(self, time):
        """
        Computes sinusoidal embeddings for given time steps.

        Args:
            time (torch.Tensor): Time steps tensor of shape [batch_size]

        Returns:
            torch.Tensor: Time embeddings of shape [batch_size, dim]
        """
        device = time.device
        half_dim = self.dim // 2
        embeddings = torch.log(torch.tensor(10000.0, device=device)) / (half_dim - 1)
        embeddings = torch.exp(torch.arange(half_dim, device=device) * -embeddings)
        embeddings = time[:, None] * embeddings[None, :]
        embeddings = torch.cat((embeddings.sin(), embeddings.cos()), dim=-1)
        return embeddings

```

```

# Helps the model understand which number/image to draw (class conditioning)
class EmbedBlock(nn.Module):
    """
    Creates embeddings for class conditioning in diffusion models.

    This module transforms a one-hot or index representation of a class
    into a rich embedding that can be added to feature maps.

    Args:
        input_dim (int): Number of classes (input dimension)
        emb_dim (int): Output embedding dimension
    """
    def __init__(self, input_dim, emb_dim):
        super().__init__()
        self.input_dim = input_dim

        # Linear -> GELU -> Linear -> Unflatten to [B, emb_dim, 1, 1]
        self.model = nn.Sequential(
            nn.Linear(input_dim, emb_dim),
            nn.GELU(),
            nn.Linear(emb_dim, emb_dim),
            nn.Unflatten(1, (emb_dim, 1, 1)) # For broadcasting over feature maps
        )

    def forward(self, x):
        """
        Computes class embeddings for the given class indices.

```

```

Args:
    x (torch.Tensor): Class indices [B] or one-hot encodings [B, input_dim]

Returns:
    torch.Tensor: Class embeddings of shape [B, emb_dim, 1, 1]
"""

# Convert integer class labels → one-hot if needed
if x.ndim == 1 or x.shape[-1] != self.input_dim:
    x = F.one_hot(x, num_classes=self.input_dim).float()

return self.model(x)

```

```

# Main U-Net model that puts everything together
class UNet(nn.Module):
    """
    U-Net architecture for diffusion models with time and class conditioning.

    This architecture follows the standard U-Net design with:
    1. Downsampling path that reduces spatial dimensions
    2. Middle processing blocks
    3. Upsampling path that reconstructs spatial dimensions
    4. Skip connections between symmetric layers

    The model is conditioned on:
    - Time step (where we are in the diffusion process)
    - Class labels (what we want to generate)

    Args:
        T (int): Number of diffusion time steps
        img_ch (int): Number of image channels
        img_size (int): Size of input images
        down_chs (list): Channel dimensions for each level of U-Net
        t_embed_dim (int): Dimension for time embeddings
        c_embed_dim (int): Dimension for class embeddings
    """

    def __init__(self, T, img_ch, img_size, down_chs, t_embed_dim, c_embed_dim):
        super().__init__()

        # Your code to create the time embedding
        # Hint: Use SinusoidalPositionEmbedBlock, nn.Linear, and nn.GELU in sequence
        self.time_embed = nn.Sequential(
            SinusoidalPositionEmbedBlock(t_embed_dim), # ✅ FIXED: only 1 argument
            nn.Linear(t_embed_dim, t_embed_dim),
            nn.GELU(),
            nn.Linear(t_embed_dim, down_chs[-1]) # Linear layer outputs to match middle block channels
        )

        # Your code to create the class embedding
        # Hint: Use the EmbedBlock class you defined earlier
        # Corrected arguments: input_dim should be N_CLASSES, emb_dim should match middle block channels
        self.class_embed = EmbedBlock(input_dim=N_CLASSES, emb_dim=down_chs[-1]) # Corrected arguments

        # Your code to create the initial convolution
        # Hint: Use GELUConvBlock to process the input image
        self.init_conv = GELUConvBlock(img_ch, down_chs[0], group_size=8)

        # Your code to create the downsampling path
        # Hint: Use nn.ModuleList with DownBlock for each level
        self.downs = nn.ModuleList()
        for i in range(len(down_chs) - 1):
            self.downs.append(DownBlock(down_chs[i], down_chs[i+1], group_size=8))

        # Your code to create the middle blocks
        # Hint: Use GELUConvBlock twice to process features at lowest resolution
        self.mid = nn.Sequential(
            GELUConvBlock(down_chs[-1], down_chs[-1], group_size=8),
            GELUConvBlock(down_chs[-1], down_chs[-1], group_size=8)
        )

        # Your code to create the upsampling path

```

```

# Hint: Use nn.ModuleList with UpBlock for each level (in reverse order)
self.ups = nn.ModuleList()
# Determine the actual channel sizes of the skips in the order they are used in the upsampling loop
# Skips are [init_conv_out, DB1_out, DB2_out, ...]
# Skips used in loop are reversed(skips[:-1]) = [DB_last-1_out, ..., init_conv_out]
# The channels of the skips provided in the loop are:
# down_chs[len(down_chs)-2], down_chs[len(down_chs)-3], ..., down_chs[0]
actual_skip_channels_in_loop_order = [down_chs[i] for i in reversed(range(len(down_chs)-1))]

for i, (in_c, out_c, skip_c) in enumerate(zip(reversed(down_chs[1:]), reversed(down_chs[:-1]), actual_skip_channels_in_loop_order)):
    self.ups.append(UpBlock(in_c, skip_c, out_c, group_size=8))

# Your code to create the final convolution
# Hint: Use nn.Conv2d to project back to the original image channels
self.final_conv = nn.Conv2d(down_chs[0], img_ch, kernel_size=3, padding=1) # Corrected kernel_size and added padding

print(f"Created UNet with {len(down_chs)} scale levels")
print(f"Channel dimensions: {down_chs}")

def forward(self, x, t, c, c_mask):
    """
    Forward pass through the UNet.

    Args:
        x (torch.Tensor): Input noisy image [B, img_ch, H, W]
        t (torch.Tensor): Diffusion time steps [B]
        c (torch.Tensor): Class labels [B]
        c_mask (torch.Tensor): Mask for conditional generation [B, 1]

    Returns:
        torch.Tensor: Predicted noise in the input image [B, img_ch, H, W]
    """
    # Your code for the time embedding
    # Hint: Process the time steps through the time embedding module
    t_embed = self.time_embed(t)

    # Your code for the class embedding
    # Hint: Process the class labels through the class embedding module
    c_embed = self.class_embed(c)
    # if c_mask is not None: # Removed this conditional as c_mask is always passed
    c_embed = c_embed * c_mask[:, :, None, None] # Apply mask for unconditional generation

    # Your code for the initial feature extraction
    # Hint: Apply initial convolution to the input
    h = self.init_conv(x)
    skips = [h] # Store initial features as the first skip

    # Your code for the downsampling path and skip connections
    # Hint: Process the features through each downsampling block
    # and store the outputs for skip connections
    for down in self.downs:
        h = down(h)
        skips.append(h)

    # Your code for the middle processing and conditioning
    # Hint: Process features through middle blocks, then add time and class embeddings
    h = self.mid(h)

    # Expand time and class embeddings to match spatial dimensions
    # Time embedding is already expanded to match middle block channels in __init__
    # Class embedding is already expanded to match middle block channels and spatial 1x1 in __init__
    h = h + t_embed[:, :, None, None] + c_embed

    # Your code for the upsampling path with skip connections
    # Hint: Process features through each upsampling block,
    # combining with corresponding skip connections
    # The skips are [init_conv (32), DB1 (64), DB2 (128)]
    # We need to use reversed(skips[:-1]) = [DB1 (64), init_conv (32)]
    for up, skip in zip(self.ups, reversed(skips[:-1])): # Corrected skips used
        h = up(h, skip)

```

```
# Final projection
return self.final_conv(h)
```

✓ Step 4: Setting Up The Diffusion Process

Now we'll create the process of adding and removing noise from images. Think of it like:

1. Adding fog: Slowly making the image more and more blurry until you can't see it
2. Removing fog: Teaching the AI to gradually make the image clearer
3. Controlling the process: Making sure we can generate specific numbers we want

```
# Set up the noise schedule
n_steps = 100 # How many steps to go from clear image to noise
beta_start = 0.0001 # Starting noise level (small)
beta_end = 0.02 # Ending noise level (larger)

# Create schedule of gradually increasing noise levels
beta = torch.linspace(beta_start, beta_end, n_steps).to(device)

# Calculate important values used in diffusion equations
alpha = 1 - beta # Portion of original image to keep at each step
alpha_bar = torch.cumprod(alpha, dim=0) # Cumulative product of alphas
sqrt_alpha_bar = torch.sqrt(alpha_bar) # For scaling the original image
sqrt_one_minus_alpha_bar = torch.sqrt(1 - alpha_bar) # For scaling the noise
```

```
# Function to add noise to images (forward diffusion process)
def add_noise(x_0, t):
    """
    Add noise to images according to the forward diffusion process.

    The formula is:  $x_t = \sqrt{\alpha_{bar\_t}} * x_0 + \sqrt{1-\alpha_{bar\_t}} * \epsilon$ 
    where  $\epsilon$  is random noise and  $\alpha_{bar\_t}$  is the cumulative product of  $(1-\beta)$ .
    """

    Args:
        x_0 (torch.Tensor): Original clean image [B, C, H, W]
        t (torch.Tensor): Timestep indices indicating noise level [B]

    Returns:
        tuple: (noisy_image, noise_added)
            - noisy_image is the image with noise added
            - noise_added is the actual noise that was added (for training)
    """
    # Create random Gaussian noise with same shape as image
    noise = torch.randn_like(x_0)

    # Get noise schedule values for the specified timesteps
    # Reshape to allow broadcasting with image dimensions
    sqrt_alpha_bar_t = sqrt_alpha_bar[t].reshape(-1, 1, 1, 1)
    sqrt_one_minus_alpha_bar_t = sqrt_one_minus_alpha_bar[t].reshape(-1, 1, 1, 1)

    # Apply the forward diffusion equation:
    # Mixture of original image (scaled down) and noise (scaled up)
    x_t = sqrt_alpha_bar_t * x_0 + sqrt_one_minus_alpha_bar_t * noise

    return x_t, noise
```

```
# Function to remove noise from images (reverse diffusion process)
@torch.no_grad() # Don't track gradients during sampling (inference only)
def remove_noise(x_t, t, model, c, c_mask):
    """
    Remove noise from images using the learned reverse diffusion process.

    This implements a single step of the reverse diffusion sampling process.
    The model predicts the noise in the image, which we then use to partially
    denoise the image.
    """

    pass
```

```

Args:
    x_t (torch.Tensor): Noisy image at timestep t [B, C, H, W]
    t (torch.Tensor): Current timestep indices [B]
    model (nn.Module): U-Net model that predicts noise
    c (torch.Tensor): Class conditioning (what digit to generate) [B, C]
    c_mask (torch.Tensor): Mask for conditional generation [B, 1]

Returns:
    torch.Tensor: Less noisy image for the next timestep [B, C, H, W]
"""
# Predict the noise in the image using our model
predicted_noise = model(x_t, t, c, c_mask)

# Get noise schedule values for the current timestep
alpha_t = alpha[t].reshape(-1, 1, 1, 1)
alpha_bar_t = alpha_bar[t].reshape(-1, 1, 1, 1)
beta_t = beta[t].reshape(-1, 1, 1, 1)
sqrt_one_minus_alpha_bar_t = sqrt_one_minus_alpha_bar[t].reshape(-1, 1, 1, 1)

# Special case: if we're at the first timestep (t=0), we're done
if t[0] == 0:
    return x_t
else:
    # Calculate the mean of the denoised distribution
    # This is derived from Bayes' rule and the diffusion process equations
    mean = (1 / torch.sqrt(alpha_t)) * (
        x_t - (beta_t / sqrt_one_minus_alpha_bar_t) * predicted_noise
    )

    # Add a small amount of random noise (variance depends on timestep)
    # This helps prevent the generation from becoming too deterministic
    noise = torch.randn_like(x_t)

    # Return the partially denoised image with a bit of new random noise
return mean + torch.sqrt(beta_t) * noise

```

```

# Visualization function to show how noise progressively affects images
def show_noise_progression(image, num_steps=5):
    """
    Visualize how an image gets progressively noisier in the diffusion process.

Args:
    image (torch.Tensor): Original clean image [C, H, W]
    num_steps (int): Number of noise levels to show
"""
    plt.figure(figsize=(15, 3))

    # Show original image
    plt.subplot(1, num_steps, 1)
    if IMG_CH == 1: # Grayscale image
        plt.imshow(image[0].cpu(), cmap='gray')
    else: # Color image
        img = image.permute(1, 2, 0).cpu() # Change from [C,H,W] to [H,W,C]
        if img.min() < 0: # If normalized between -1 and 1
            img = (img + 1) / 2 # Rescale to [0,1] for display
        plt.imshow(img)
    plt.title('Original')
    plt.axis('off')

    # Show progressively noisier versions
    for i in range(1, num_steps):
        # Calculate timestep index based on percentage through the process
        t_idx = int((i/num_steps) * n_steps)
        t = torch.tensor([t_idx]).to(device)

        # Add noise corresponding to timestep t
        noisy_image, _ = add_noise(image.unsqueeze(0), t)


        # Display the noisy image
        plt.subplot(1, num_steps, i+1)
        if IMG_CH == 1:
            plt.imshow(noisy_image[0][0].cpu(), cmap='gray')
        else:

```

```



```



```
'\n# Try a non-linear noise schedule\nbeta_alt = torch.linspace(beta_start, beta_end, n_steps)**2\nalpha_alt = 1 - beta_alt\nalpha_bar_alt = torch.cumprod(alpha_alt, dim=0)\n# How would this affect the diffusion process?\n'
```

▼ Step 5: Training Our Model

Now we'll teach our AI to generate images. This process:

1. Takes a clear image
2. Adds random noise to it
3. Asks our AI to predict what noise was added
4. Helps our AI learn from its mistakes

This will take a while, but we'll see progress as it learns!

```

=====
# MODEL CREATION, OPTIMIZER, AND VALIDATION
=====

# Create our model and move it to GPU if available
model = UNet(
    T=n_steps,                      # Number of diffusion time steps
    img_ch=IMG_CH,                  # Number of channels in our images (1 for grayscale, 3 for RGB)
    img_size=IMG_SIZE,               # Size of input images (28 for MNIST, 32 for CIFAR-10)
    down_chs=(32, 64, 128),          # Channel dimensions for each downsampling level
    t_embed_dim=8,                  # Dimension for time step embeddings
    c_embed_dim=N_CLASSES           # Number of classes for conditioning
).to(device)

# Print model summary
print(f"\n{'='*50}")
print(f"MODEL ARCHITECTURE SUMMARY")
print(f"{'='*50}")
print(f"Input resolution: {IMG_SIZE}x{IMG_SIZE}")
print(f"Input channels: {IMG_CH}")
print(f"Time steps: {n_steps}")

```

```

print(f"Condition classes: {N_CLASSES}")
print(f"GPU acceleration: {'Yes' if device.type == 'cuda' else 'No'}")

#-----
# Validate model parameters and memory requirements
#-----

def count_parameters(model):
    return sum(p.numel() for p in model.parameters() if p.requires_grad)

num_params = count_parameters(model)
print(f"Total trainable parameters: {num_params:,}")

# Rough GPU memory estimate: assume 4 bytes per parameter + activations (~2-3x)
approx_mem_mb = num_params * 4 / (1024**2) * 3
print(f"Approximate memory usage: {approx_mem_mb:.1f} MB")

#-----
# Verify data ranges and integrity
#-----

def check_data_ranges(loader):
    for images, labels in loader:
        print(f"Sample batch shape: {images.shape}")
        print(f"Min pixel value: {images.min().item():.3f}")
        print(f"Max pixel value: {images.max().item():.3f}")
        print(f"Labels: {labels[:10]}")
        break # Only check the first batch

check_data_ranges(train_dataloader)

#-----
# Set up the optimizer
#-----

initial_lr = 0.001 # Starting learning rate
weight_decay = 1e-5 # L2 regularization

optimizer = torch.optim.Adam(
    model.parameters(),
    lr=initial_lr,
    weight_decay=weight_decay
)

# Learning rate scheduler to reduce LR when validation loss plateaus
# Note: Removed verbose argument for compatibility
scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(
    optimizer,
    mode='min',
    factor=0.5,
    patience=5,
    min_lr=1e-6
)

#-----
# STUDENT EXPERIMENT:
# Try different channel configurations and see how they affect:
# 1. Model size (parameter count)
# 2. Training time
# 3. Generated image quality
# Suggestions:
# - Smaller: down_chs=(16, 32, 64)
# - Larger: down_chs=(64, 128, 256, 512)
#-----

```

```

Created DownBlock: in_chs=32, out_chs=64, spatial_reduction=2x
Created DownBlock: in_chs=64, out_chs=128, spatial_reduction=2x
Created UpBlock: in_chs=128, skip_chs=64, out_chs=64, spatial_increase=2x
Created UpBlock: in_chs=64, skip_chs=32, out_chs=32, spatial_increase=2x
Created UNet with 3 scale levels
Channel dimensions: (32, 64, 128)

```

```

=====
MODEL ARCHITECTURE SUMMARY
=====
Input resolution: 32x32
Input channels: 1
Time steps: 100

```

```
Condition classes: 10
GPU acceleration: Yes
Total trainable parameters: 957,513
Approximate memory usage: 11.0 MB
Sample batch shape: torch.Size([64, 1, 28, 28])
Min pixel value: -1.000
Max pixel value: 1.000
Labels: tensor([1, 9, 2, 4, 4, 1, 8, 9, 1, 9])
```

```
# Define helper functions needed for training and evaluation
def validate_model_parameters(model):
    """
    Counts model parameters and estimates memory usage.
    """
    total_params = sum(p.numel() for p in model.parameters())
    trainable_params = sum(p.numel() for p in model.parameters() if p.requires_grad)

    print(f"Total parameters: {total_params},")
    print(f"Trainable parameters: {trainable_params},")

    # Estimate memory requirements (very approximate)
    param_memory = total_params * 4 / (1024 ** 2) # MB for params (float32)
    grad_memory = trainable_params * 4 / (1024 ** 2) # MB for gradients
    buffer_memory = param_memory * 2 # Optimizer state, forward activations, etc.

    print(f"Estimated GPU memory usage: {param_memory + grad_memory + buffer_memory:.1f} MB")

# Define helper functions for verifying data ranges
def verify_data_range(dataloader, name="Dataset"):
    """
    Verifies the range and integrity of the data.
    """
    batch = next(iter(dataloader))[0]
    print(f"\n{name} range check:")
    print(f"Shape: {batch.shape}")
    print(f"Data type: {batch.dtype}")
    print(f"Min value: {batch.min().item():.2f}")
    print(f"Max value: {batch.max().item():.2f}")
    print(f"Contains NaN: {torch.isnan(batch).any().item()}")
    print(f"Contains Inf: {torch.isinf(batch).any().item()}")

# Define helper functions for generating samples during training
def generate_samples(model, n_samples=10):
    """
    Generates sample images using the model for visualization during training.
    """
    model.eval()
    with torch.no_grad():
        # Generate digits 0-9 for visualization
        samples = []
        for digit in range(min(n_samples, 10)):
            # Start with random noise
            x = torch.randn(1, IMG_CH, IMG_SIZE, IMG_SIZE).to(device)

            # Set up conditioning for the digit
            c = torch.tensor([digit]).to(device)
            c_one_hot = F.one_hot(c, N_CLASSES).float().to(device)
            c_mask = torch.ones_like(c.unsqueeze(-1)).to(device)

            # Remove noise step by step
            for t in range(n_steps-1, -1, -1):
                t_batch = torch.full((1,), t).to(device)
                x = remove_noise(x, t_batch, model, c_one_hot, c_mask)

            samples.append(x)

        # Combine samples and display
        samples = torch.cat(samples, dim=0)
        grid = make_grid(samples, nrow=min(n_samples, 5), normalize=True)

        plt.figure(figsize=(10, 4))

        # Display based on channel configuration
        if IMG_CH == 1:
            plt.imshow(grid[0].cpu(), cmap='gray')
```

```

        else:
            plt.imshow(grid.permute(1, 2, 0).cpu())

            plt.axis('off')
            plt.title('Generated Samples')
            plt.show()

# Define helper functions for safely saving models
def safe_save_model(model, path, optimizer=None, epoch=None, best_loss=None):
    """
    Safely saves model with error handling and backup.
    """
    try:
        # Create a dictionary with all the elements to save
        save_dict = {
            'model_state_dict': model.state_dict(),
        }

        # Add optional elements if provided
        if optimizer is not None:
            save_dict['optimizer_state_dict'] = optimizer.state_dict()
        if epoch is not None:
            save_dict['epoch'] = epoch
        if best_loss is not None:
            save_dict['best_loss'] = best_loss

        # Create a backup of previous checkpoint if it exists
        if os.path.exists(path):
            backup_path = path + '.backup'
            try:
                os.replace(path, backup_path)
                print(f"Created backup at {backup_path}")
            except Exception as e:
                print(f"Warning: Could not create backup - {e}")

        # Save the new checkpoint
        torch.save(save_dict, path)
        print(f"Model successfully saved to {path}")

    except Exception as e:
        print(f"Error saving model: {e}")
        print("Attempting emergency save...")

    try:
        emergency_path = path + '.emergency'
        torch.save(model.state_dict(), emergency_path)
        print(f"Emergency save successful: {emergency_path}")
    except:
        print("Emergency save failed. Could not save model.")

```

```

# Implementation of the training step function
def train_step(x, c):
    """
    Performs a single training step for the diffusion model.

    This function:
    1. Prepares class conditioning
    2. Samples random timesteps for each image
    3. Adds corresponding noise to the images
    4. Asks the model to predict the noise
    5. Calculates the loss between predicted and actual noise

    Args:
        x (torch.Tensor): Batch of clean images [batch_size, channels, height, width]
        c (torch.Tensor): Batch of class labels [batch_size]

    Returns:
        torch.Tensor: Mean squared error loss value
    """
    # Convert number labels to one-hot encoding for class conditioning
    # Example: Label 3 -> [0, 0, 0, 1, 0, 0, 0, 0, 0] for MNIST
    c_one_hot = F.one_hot(c, N_CLASSES).float().to(device)

    # Create conditioning mask (all ones for standard training)

```

```

# This would be used for classifier-free guidance if implemented
c_mask = torch.ones_like(c.unsqueeze(-1)).to(device)

# Pick random timesteps for each image in the batch
# Different timesteps allow the model to learn the entire diffusion process
t = torch.randint(0, n_steps, (x.shape[0],)).to(device)

# Add noise to images according to the forward diffusion process
# This simulates images at different stages of the diffusion process
# Hint: Use the add_noise function you defined earlier

# Enter your code here: # Removed placeholder comment
x_t, noise = add_noise(x, t) # Added call to add_noise

# The model tries to predict the exact noise that was added
# This is the core learning objective of diffusion models
predicted_noise = model(x_t, t, c_one_hot, c_mask)

# Calculate loss: how accurately did the model predict the noise?
# MSE loss works well for image-based diffusion models
# Hint: Use F.mse_loss to compare predicted and actual noise

# Enter your code here: # Removed placeholder comment
loss = F.mse_loss(predicted_noise, noise) # Added loss calculation

```

```
return loss
```

```

# Implementation of the main training loop
# Training configuration
early_stopping_patience = 10 # Number of epochs without improvement before stopping
gradient_clip_value = 1.0 # Maximum gradient norm for stability
display_frequency = 100 # How often to show progress (in steps)
generate_frequency = 500 # How often to generate samples (in steps)

# Progress tracking variables
best_loss = float('inf')
train_losses = []
val_losses = []
no_improve_epochs = 0

# Training loop
print("\n" + "="*50)
print("STARTING TRAINING")
print("="*50)

# Wrap the training loop in a try-except block for better error handling
try:
    model.train() # Set model to training mode
    for epoch in range(EPOCHS):
        print(f"\nEpoch {epoch+1}/{EPOCHS}")
        print("-" * 20)

        # Training phase
        model.train()
        epoch_losses = []

        # Process each batch
        for step, (images, labels) in enumerate(train_dataloader): # Iterate over training data
            images = images.to(device)
            labels = labels.to(device)

            # Training step
            optimizer.zero_grad() # Reset gradients
            loss = train_step(images, labels) # Compute loss
            loss.backward() # Backpropagate

            # Add gradient clipping for stability
            torch.nn.utils.clip_grad_norm_(model.parameters(), max_norm=gradient_clip_value)

            optimizer.step() # Update weights
            epoch_losses.append(loss.item())

        # Show progress at regular intervals

```

```

if step % display_frequency == 0:
    print(f" Step {step}/{len(train_dataloader)}, Loss: {loss.item():.4f}")

    # Generate samples less frequently to save time
    if step % generate_frequency == 0 and step > 0:
        print(" Generating samples...")
        generate_samples(model, n_samples=5)

# End of epoch - calculate average training loss
avg_train_loss = sum(epoch_losses) / len(epoch_losses)
train_losses.append(avg_train_loss)
print(f"\nTraining - Epoch {epoch+1} average loss: {avg_train_loss:.4f}")

# Validation phase
model.eval() # Set model to evaluation mode
val_epoch_losses = []
print("Running validation...")

with torch.no_grad(): # Disable gradients for validation
    for val_images, val_labels in val_dataloader:
        val_images = val_images.to(device)
        val_labels = val_labels.to(device)

        # Calculate validation loss
        val_loss = train_step(val_images, val_labels)
        val_epoch_losses.append(val_loss.item())

# Calculate average validation loss
avg_val_loss = sum(val_epoch_losses) / len(val_epoch_losses)
val_losses.append(avg_val_loss)
print(f"Validation - Epoch {epoch+1} average loss: {avg_val_loss:.4f}")

# Learning rate scheduling based on validation loss
scheduler.step(avg_val_loss)
current_lr = optimizer.param_groups[0]['lr']
print(f"Learning rate: {current_lr:.6f}")

# Generate samples at the end of each epoch for visual check
if epoch % 2 == 0 or epoch == EPOCHS - 1:
    print("\nGenerating samples for visual progress check...")
    generate_samples(model, n_samples=10)

# Save best model based on validation loss
if avg_val_loss < best_loss:
    best_loss = avg_val_loss
    safe_save_model(model, 'best_diffusion_model.pt', optimizer, epoch, best_loss)
    print(f"/ New best model saved! (Val Loss: {best_loss:.4f})")
    no_improve_epochs = 0
else:
    no_improve_epochs += 1
    print(f"No improvement for {no_improve_epochs}/{early_stopping_patience} epochs")

# Early stopping
if no_improve_epochs >= early_stopping_patience:
    print("\nEarly stopping triggered! No improvement in validation loss.")
    break

# Plot loss curves every few epochs
if epoch % 5 == 0 or epoch == EPOCHS - 1:
    plt.figure(figsize=(10, 5))
    plt.plot(train_losses, label='Training Loss')
    plt.plot(val_losses, label='Validation Loss')
    plt.xlabel('Epoch')
    plt.ylabel('Loss')
    plt.title('Training and Validation Loss')
    plt.legend()
    plt.grid(True)
    plt.show()

except Exception as e: # Catch errors during training
    print(f"An error occurred during training: {e}")
    import traceback
    traceback.print_exc()

```

```
# Final wrap-up
print("\n" + "="*50)
print("TRAINING COMPLETE")
print("=*50)
print(f"Best validation loss: {best_loss:.4f}")

# Generate final samples
print("Generating final samples...")
generate_samples(model, n_samples=10)

# Display final loss curves
plt.figure(figsize=(12, 5))
plt.plot(train_losses, label='Training Loss')
plt.plot(val_losses, label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training and Validation Loss')
plt.legend()
plt.grid(True)
plt.show()

# Clean up memory
torch.cuda.empty_cache()
```




```
=====
STARTING TRAINING
=====

Epoch 1/30
-----
Step 0/352, Loss: 1.1633
Step 100/352, Loss: 0.1179
Step 200/352, Loss: 0.0946
Step 300/352, Loss: 0.0957
```


Training - Epoch 1 average loss: 0.1284
 Running validation...
 Validation - Epoch 1 average loss: 0.0854
 Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0854)



Epoch 2/30

 Step 0/352, Loss: 0.0738
 Step 100/352, Loss: 0.0862
 Step 200/352, Loss: 0.0678
 Step 300/352, Loss: 0.0662

Training - Epoch 2 average loss: 0.0785
 Running validation...
 Validation - Epoch 2 average loss: 0.0761
 Learning rate: 0.001000
 Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0761)

Epoch 3/30

```
-----  
Step 0/352, Loss: 0.0685  
Step 100/352, Loss: 0.0624  
Step 200/352, Loss: 0.0802  
Step 300/352, Loss: 0.0733
```

Training - Epoch 3 average loss: 0.0724

Running validation...

Validation - Epoch 3 average loss: 0.0700

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



Created backup at best_diffusion_model.pt.backup

Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.0700)

Epoch 4/30

```
-----  
Step 0/352, Loss: 0.0672  
Step 100/352, Loss: 0.0639  
Step 200/352, Loss: 0.0697  
Step 300/352, Loss: 0.0566
```

Training - Epoch 4 average loss: 0.0690

Running validation...

Validation - Epoch 4 average loss: 0.0688

Learning rate: 0.001000

Created backup at best_diffusion_model.pt.backup

Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.0688)

Epoch 5/30

```
-----  
Step 0/352, Loss: 0.0733  
Step 100/352, Loss: 0.0604  
Step 200/352, Loss: 0.0633  
Step 300/352, Loss: 0.0658
```

Training - Epoch 5 average loss: 0.0665


Running validation...

Validation - Epoch 5 average loss: 0.0657

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples





Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0657)

Epoch 6/30

```
Step 0/352, Loss: 0.0651
Step 100/352, Loss: 0.0662
Step 200/352, Loss: 0.0574
Step 300/352, Loss: 0.0668
```

Training - Epoch 6 average loss: 0.0657


Running validation...

Validation - Epoch 6 average loss: 0.0668

Learning rate: 0.001000

No improvement for 1/10 epochs

Training and Validation Loss



Epoch 7/30

```
Step 0/352, Loss: 0.0682
Step 100/352, Loss: 0.0574
Step 200/352, Loss: 0.0714
Step 300/352, Loss: 0.0654
```

Training - Epoch 7 average loss: 0.0646

Running validation...

Validation - Epoch 7 average loss: 0.0637

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt

```
model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.0637)
```

Epoch 8/30

```
Step 0/352, Loss: 0.0628
Step 100/352, Loss: 0.0620
Step 200/352, Loss: 0.0535
Step 300/352, Loss: 0.0755
```

Training - Epoch 8 average loss: 0.0641

Running validation...

Validation - Epoch 8 average loss: 0.0630

Learning rate: 0.001000

Created backup at best_diffusion_model.pt.backup

Model successfully saved to best_diffusion_model.pt

```
✓ New best model saved! (Val Loss: 0.0630)
```

Epoch 9/30

```
Step 0/352, Loss: 0.0669
Step 100/352, Loss: 0.0575
Step 200/352, Loss: 0.0746
Step 300/352, Loss: 0.0610
```

Training - Epoch 9 average loss: 0.0634


Running validation...

Validation - Epoch 9 average loss: 0.0633

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



No improvement for 1/10 epochs

Epoch 10/30

```
Step 0/352, Loss: 0.0647
Step 100/352, Loss: 0.0572
Step 200/352, Loss: 0.0629
Step 300/352, Loss: 0.0577
```

Training - Epoch 10 average loss: 0.0627

Running validation...

Validation - Epoch 10 average loss: 0.0604

Learning rate: 0.001000

Created backup at best_diffusion_model.pt.backup

Model successfully saved to best_diffusion_model.pt

```
✓ New best model saved! (Val Loss: 0.0604)
```

Epoch 11/30

```
Step 0/352, Loss: 0.0550
Step 100/352, Loss: 0.0651
Step 200/352, Loss: 0.0581
Step 300/352, Loss: 0.0519
```

Training - Epoch 11 average loss: 0.0619


Running validation...

Validation - Epoch 11 average loss: 0.0624

Learning rate: 0.001000


Generating samples for visual progress check...

Generated Samples



No improvement for 1/10 epochs

Training and Validation Loss



Epoch 12/30

```
Step 0/352, Loss: 0.0594
Step 100/352, Loss: 0.0583
Step 200/352, Loss: 0.0546
Step 300/352, Loss: 0.0617
```

Training - Epoch 12 average loss: 0.0616
Running validation...

Validation - Epoch 12 average loss: 0.0639

Learning rate: 0.001000

No improvement for 2/10 epochs

Epoch 13/30

```
Step 0/352, Loss: 0.0656
Step 100/352, Loss: 0.0623
Step 200/352, Loss: 0.0675
Step 300/352, Loss: 0.0687
```

Training - Epoch 13 average loss: 0.0609
Running validation...

Validation - Epoch 13 average loss: 0.0608

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples





No improvement for 3/10 epochs

Epoch 14/30

```
Step 0/352, Loss: 0.0586
Step 100/352, Loss: 0.0648
Step 200/352, Loss: 0.0721
Step 300/352, Loss: 0.0543
```

Training - Epoch 14 average loss: 0.0611

Running validation...

Validation - Epoch 14 average loss: 0.0603

Learning rate: 0.001000

```
Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.0603)
```

Epoch 15/30

```
Step 0/352, Loss: 0.0493
Step 100/352, Loss: 0.0660
Step 200/352, Loss: 0.0698
Step 300/352, Loss: 0.0637
```

Training - Epoch 15 average loss: 0.0605

Running validation...

Validation - Epoch 15 average loss: 0.0593

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



Created backup at best_diffusion_model.pt.backup

Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.0593)

Epoch 16/30

```
Step 0/352, Loss: 0.0560
Step 100/352, Loss: 0.0610
Step 200/352, Loss: 0.0629
Step 300/352, Loss: 0.0651
```

Training - Epoch 16 average loss: 0.0607


Running validation...


Validation - Epoch 16 average loss: 0.0613

Learning rate: 0.001000

No improvement for 1/10 epochs

Training and Validation Loss





Epoch 17/30

```
Step 0/352, Loss: 0.0544
Step 100/352, Loss: 0.0584
Step 200/352, Loss: 0.0584
Step 300/352, Loss: 0.0504
```

Training - Epoch 17 average loss: 0.0596

Running validation...

Validation - Epoch 17 average loss: 0.0599

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



No improvement for 2/10 epochs

Epoch 18/30

```
Step 0/352, Loss: 0.0578
Step 100/352, Loss: 0.0695
Step 200/352, Loss: 0.0560
Step 300/352, Loss: 0.0595
```

Training - Epoch 18 average loss: 0.0599

Running validation...

Validation - Epoch 18 average loss: 0.0595

Learning rate: 0.001000

No improvement for 3/10 epochs

Epoch 19/30

```
Step 0/352, Loss: 0.0715
Step 100/352, Loss: 0.0620
Step 200/352, Loss: 0.0553
Step 300/352, Loss: 0.0624
```

Training - Epoch 19 average loss: 0.0599

Running validation...

Validation - Epoch 19 average loss: 0.0586

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



Created backup at best_diffusion_model.pt.backup
Model successfully saved to best_diffusion_model.pt
✓ New best model saved! (Val Loss: 0.0586)

Epoch 20/30

```
-----  
Step 0/352, Loss: 0.0565  
Step 100/352, Loss: 0.0578  
Step 200/352, Loss: 0.0611  
Step 300/352, Loss: 0.0608
```

Training - Epoch 20 average loss: 0.0591

Running validation...

Validation - Epoch 20 average loss: 0.0604

Learning rate: 0.001000

No improvement for 1/10 epochs

Epoch 21/30

```
-----  
Step 0/352, Loss: 0.0573  
Step 100/352, Loss: 0.0542  
Step 200/352, Loss: 0.0514  
Step 300/352, Loss: 0.0511
```

Training - Epoch 21 average loss: 0.0595

Running validation...

Validation - Epoch 21 average loss: 0.0613

Learning rate: 0.001000


Generating samples for visual progress check...


Generated Samples



No improvement for 2/10 epochs

Training and Validation Loss





Epoch 22/30

```
Step 0/352, Loss: 0.0592
Step 100/352, Loss: 0.0529
Step 200/352, Loss: 0.0563
Step 300/352, Loss: 0.0545
```

Training - Epoch 22 average loss: 0.0590

Running validation...

Validation - Epoch 22 average loss: 0.0582

Learning rate: 0.001000

Created backup at best_diffusion_model.pt.backup

Model successfully saved to best_diffusion_model.pt

✓ New best model saved! (Val Loss: 0.0582)

Epoch 23/30

```
Step 0/352, Loss: 0.0582
Step 100/352, Loss: 0.0552
Step 200/352, Loss: 0.0492
Step 300/352, Loss: 0.0523
```

Training - Epoch 23 average loss: 0.0582

Running validation...

Validation - Epoch 23 average loss: 0.0602

Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



No improvement for 1/10 epochs

Epoch 24/30

```
Step 0/352, Loss: 0.0496
Step 100/352, Loss: 0.0525
Step 200/352, Loss: 0.0524
Step 300/352, Loss: 0.0595
```

Training - Epoch 24 average loss: 0.0589

Running validation...

Validation - Epoch 24 average loss: 0.0589

Learning rate: 0.001000

No improvement for 2/10 epochs


Epoch 25/30

Step 0/352, Loss: 0.0582
 Step 100/352, Loss: 0.0567
 Step 200/352, Loss: 0.0645
 Step 300/352, Loss: 0.0551

Training - Epoch 25 average loss: 0.0585
 Running validation...
 Validation - Epoch 25 average loss: 0.0578
 Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples




Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0578)

Epoch 26/30

Step 0/352, Loss: 0.0534
 Step 100/352, Loss: 0.0524
 Step 200/352, Loss: 0.0621
 Step 300/352, Loss: 0.0699

Training - Epoch 26 average loss: 0.0584
 Running validation...
 Validation - Epoch 26 average loss: 0.0575
 Learning rate: 0.001000
 Created backup at best_diffusion_model.pt.backup
 Model successfully saved to best_diffusion_model.pt
 ✓ New best model saved! (Val Loss: 0.0575)

Training and Validation Loss



Epoch 27/30

Step 0/352, Loss: 0.0588
 Step 100/352, Loss: 0.0587
 Step 200/352, Loss: 0.0673

```
Step 200/352, Loss: 0.0075  
Step 300/352, Loss: 0.0571
```

Training - Epoch 27 average loss: 0.0584
Running validation...
Validation - Epoch 27 average loss: 0.0600
Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



No improvement for 1/10 epochs

Epoch 28/30

```
-----  
Step 0/352, Loss: 0.0641  
Step 100/352, Loss: 0.0556  
Step 200/352, Loss: 0.0483  
Step 300/352, Loss: 0.0572
```

Training - Epoch 28 average loss: 0.0581
Running validation...
Validation - Epoch 28 average loss: 0.0577
Learning rate: 0.001000
No improvement for 2/10 epochs

Epoch 29/30

```
-----  
Step 0/352, Loss: 0.0664  
Step 100/352, Loss: 0.0576  
Step 200/352, Loss: 0.0595  
Step 300/352, Loss: 0.0575
```

Training - Epoch 29 average loss: 0.0583
Running validation...
Validation - Epoch 29 average loss: 0.0596
Learning rate: 0.001000

Generating samples for visual progress check...

Generated Samples



No improvement for 3/10 epochs

Epoch 30/30

```
-----  
Step 0/352, Loss: 0.0549  
Step 100/352, Loss: 0.0583
```

```

# Plot training progress
plt.figure(figsize=(12, 5))

# Plot training and validation losses for comparison
plt.plot(train_losses, label='Training Loss')
if len(val_losses) > 0: # Only plot validation if it exists
    plt.plot(val_losses, label='Validation Loss')

# Improve the plot with better labels and styling
plt.title('Diffusion Model Training Progress')
plt.xlabel('Epoch')
plt.ylabel('Loss (MSE)')
plt.legend()
plt.grid(True)

# Add annotations for key points
if len(train_losses) > 1:
    min_train_idx = train_losses.index(min(train_losses))
    plt.annotate(f'Min: {min(train_losses):.4f}',
                 xy=(min_train_idx, min(train_losses)),
                 xytext=(min_train_idx, min(train_losses)*1.2),
                 arrowprops=dict(facecolor='black', shrink=0.05),
                 fontsize=9)

# Add validation min point if available
if len(val_losses) > 1:
    min_val_idx = val_losses.index(min(val_losses))
    plt.annotate(f'Min: {min(val_losses):.4f}',
                 xy=(min_val_idx, min(val_losses)),
                 xytext=(min_val_idx, min(val_losses)*0.8),
                 arrowprops=dict(facecolor='black', shrink=0.05),
                 fontsize=9)

# Set y-axis to start from 0 or slightly lower than min value
plt.ylim(bottom=max(0, min(min(train_losses)) if train_losses else float('inf'),
                      min(val_losses) if val_losses else float('inf'))*0.9))


plt.tight_layout()
plt.show()


# Add statistics summary for students to analyze
print("\nTraining Statistics:")
print("-" * 30)
if train_losses:
    print(f"Starting training loss: {train_losses[0]:.4f}")
    print(f"Final training loss: {train_losses[-1]:.4f}")
    print(f"Best training loss: {min(train_losses):.4f}")
    print(f"Training loss improvement: {((train_losses[0] - min(train_losses)) / train_losses[0] * 100):.1f}%")

if val_losses:
    print("\nValidation Statistics:")
    print("-" * 30)
    print(f"Starting validation loss: {val_losses[0]:.4f}")
    print(f"Final validation loss: {val_losses[-1]:.4f}")
    print(f"Best validation loss: {min(val_losses):.4f}")

# STUDENT EXERCISE:
# 1. Try modifying this plot to show a smoothed version of the losses
# 2. Create a second plot showing the ratio of validation to training loss
#     (which can indicate overfitting when the ratio increases)

```





▼ Step 6: Generating New Images

Now that our model is trained, let's generate some new images! We can:

1. Generate specific numbers
2. Generate multiple versions of each number
3. See how the generation process works step by step

```
def generate_number(model, number, n_samples=4):
    """
    Generate multiple versions of a specific number using the diffusion model.

    Args:
        model (nn.Module): The trained diffusion model
        number (int): The digit to generate (0-9)
        n_samples (int): Number of variations to generate

    Returns:
        torch.Tensor: Generated images of shape [n_samples, IMG_CH, IMG_SIZE, IMG_SIZE]
    """
    model.eval() # Set model to evaluation mode
    with torch.no_grad(): # No need for gradients during generation
        # Start with random noise
        samples = torch.randn(n_samples, IMG_CH, IMG_SIZE, IMG_SIZE).to(device)

        # Set up the number we want to generate
        c = torch.full((n_samples,), number).to(device)
        c_one_hot = F.one_hot(c, N_CLASSES).float().to(device)
        # Correctly sized conditioning mask
        c_mask = torch.ones_like(c.unsqueeze(-1)).to(device)

        # Display progress information
        print(f"Generating {n_samples} versions of number {number}...")

        # Remove noise step by step
        for t in range(n_steps-1, -1, -1):
            t_batch = torch.full((n_samples,), t).to(device)
```

```

samples = remove_noise(samples, t_batch, model, c_one_hot, c_mask)

    # Optional: Display occasional progress updates
    if t % (n_steps // 5) == 0:
        print(f" Denoising step {n_steps-1-t}/{n_steps-1} completed")

return samples

# Generate 4 versions of each number
plt.figure(figsize=(20, 10))
for i in range(10):
    # Generate samples for current digit
    samples = generate_number(model, i, n_samples=4)

    # Display each sample
    for j in range(4):
        # Use 2 rows, 10 digits per row, 4 samples per digit
        # i//5 determines the row (0 or 1)
        # i%5 determines the position in the row (0-4)
        # j is the sample index within each digit (0-3)
        plt.subplot(5, 8, (i%5)*8 + (i//5)*4 + j + 1)

        # Display the image correctly based on channel configuration
        if IMG_CH == 1: # Grayscale
            plt.imshow(samples[j][0].cpu(), cmap='gray')
        else: # Color image
            img = samples[j].permute(1, 2, 0).cpu()
            # Rescale from [-1, 1] to [0, 1] if needed
            if img.min() < 0:
                img = (img + 1) / 2
            plt.imshow(img)

        plt.title(f'Digit {i}')
        plt.axis('off')

plt.tight_layout()
plt.show()

# STUDENT ACTIVITY: Try generating the same digit with different noise seeds
# This shows the variety of styles the model can produce
print("\nSTUDENT ACTIVITY: Generating numbers with different noise seeds")

# Helper function to generate with seed
def generate_with_seed(number, seed_value=42, n_samples=10):
    torch.manual_seed(seed_value)
    return generate_number(model, number, n_samples)

# Pick a image and show many variations
# Hint select a image e.g. dog # Change this to any other in the dataset or subset you chose
# Hint 2 use variations = generate_with_seed
# Hint 3 use plt.figure and plt.imshow to display the variations

# Enter your code here:

# Choose the digit (or class) to visualize
digit_to_test = 5 # 🚧 Try changing this number (0-9)
n_variations = 6 # Number of variations to show

plt.figure(figsize=(12, 4))

# Try different random seeds to see how the style changes
seeds = [0, 21, 42, 84, 123, 999]
for idx, seed in enumerate(seeds):
    torch.manual_seed(seed) # Set the random seed
    variations = generate_number(model, digit_to_test, n_samples=1)

    plt.subplot(1, n_variations, idx + 1)
    if IMG_CH == 1:
        plt.imshow(variations[0][0].cpu(), cmap='gray')
    else:
        img = variations[0].permute(1, 2, 0).cpu()
        if img.min() < 0:
            img = (img + 1) / 2 # Rescale to [0,1]

```

```
plt.imshow(img)

plt.title(f"Seed {seed}")
plt.axis('off')


plt.suptitle(f"Variations of Digit {digit_to_test} with Different Seeds", fontsize=14)
plt.tight_layout()
plt.show()
```



```

Generating 4 versions of number 0...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 1...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 2...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 3...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 4...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 5...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 6...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 7...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 8...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed
Generating 4 versions of number 9...
Denoising step 19/99 completed
Denoising step 39/99 completed
Denoising step 59/99 completed
Denoising step 79/99 completed
Denoising step 99/99 completed

```



▼ Step 7: Watching the Generation Process

Let's see how our model turns random noise into clear images step by step. This helps us understand how the diffusion process works!

```

def visualize_generation_steps(model, number, n_preview_steps=10):
    """
    Show how an image evolves from noise to a clear number
    """
    model.eval()
    with torch.no_grad():
        # Start with random noise
        x = torch.randn(1, IMG_CH, IMG_SIZE, IMG_SIZE).to(device)

        # Set up which number to generate
        c = torch.tensor([number]).to(device)
        c_one_hot = F.one_hot(c, N_CLASSES).float().to(device)

```

```
# Fix: make mask have shape [1, 1], not [1, 10]
c_mask = torch.ones((1, 1)).to(device)

# Calculate which steps to show
steps_to_show = torch.linspace(n_steps-1, 0, n_preview_steps).long()


# Store images for visualization
images = [x[0].cpu()]

# Remove noise step by step
for t in range(n_steps-1, -1, -1):
    t_batch = torch.full((1,), t).to(device)
    x = remove_noise(x, t_batch, model, c_one_hot, c_mask)

    if t in steps_to_show:
        images.append(x[0].cpu())

# Show the progression
plt.figure(figsize=(20, 3))
for i, img in enumerate(images):
    plt.subplot(1, len(images), i+1)
    if IMG_CH == 1:
        plt.imshow(img[0], cmap='gray')
    else:
        img = img.permute(1, 2, 0)
        if img.min() < 0:
            img = (img + 1) / 2
        plt.imshow(img)
    step = n_steps if i == 0 else steps_to_show[i-1]
    plt.title(f'Step {step}')
    plt.axis('off')
plt.show()

# Show generation process for a few numbers
for number in [0, 3, 7]:
    print(f'\nGenerating number {number}:')
    visualize_generation_steps(model, number)
```



▼ Step 8: Adding CLIP Evaluation

[CLIP](#) is a powerful AI model that can understand both images and text. We'll use it to:

1. Evaluate how realistic our generated images are
2. Score how well they match their intended numbers
3. Help guide the generation process towards better quality

```

## Step 8: Adding CLIP Evaluation

# CLIP (Contrastive Language-Image Pre-training) is a powerful model by OpenAI that connects text and images.
# We'll use it to evaluate how recognizable our generated digits are by measuring how strongly
# the CLIP model associates our generated images with text descriptions like "an image of the digit 7".

# First, we need to install CLIP and its dependencies
print("Setting up CLIP (Contrastive Language-Image Pre-training) model...")

# Track installation status
clip_available = False

try:
    # Install dependencies first - these help CLIP process text and images
    print("Installing CLIP dependencies...")
    !pip install -q ftfy regex tqdm

    # Install CLIP from GitHub
    print("Installing CLIP from GitHub repository...")
    !pip install -q git+https://github.com/openai/CLIP.git

    # Import and verify CLIP is working
    print("Importing CLIP...")
    import clip

    # Test that CLIP is functioning
    models = clip.available_models()
    print(f"/ CLIP installation successful! Available models: {models}")
    clip_available = True

except ImportError:
    print("✖ Error importing CLIP. Installation might have failed.")
    print("Try manually running: !pip install git+https://github.com/openai/CLIP.git")
    print("If you're in a Colab notebook, try restarting the runtime after installation.")

except Exception as e:
    print(f"✖ Error during CLIP setup: {e}")
    print("Some CLIP functionality may not work correctly.")

# Provide guidance based on installation result
if clip_available:
    print("\nCLIP is now available for evaluating your generated images!")
else:
    print("\nWARNING: CLIP installation failed. We'll skip the CLIP evaluation parts.")

# Import necessary libraries
import functools
import torch.nn.functional as F

Setting up CLIP (Contrastive Language-Image Pre-training) model...
Installing CLIP dependencies... ━━━━━━━━ 44.8/44.8 kB 3.1 MB/s eta 0:00:00
Installing CLIP from GitHub repository...
  Preparing metadata (setup.py) ... done
  Building wheel for clip (setup.py) ... done
Importing CLIP...
✓ CLIP installation successful! Available models: ['RN50', 'RN101', 'RN50x4', 'RN50x16', 'RN50x64', 'ViT-B/32', 'ViT-B/16',
CLIP is now available for evaluating your generated images!

```

Below we are creating a helper function to manage GPU memory when using CLIP. CLIP can be memory-intensive, so this will help prevent out-of-memory errors:

```

# Memory management decorator to prevent GPU OOM errors
def manage_gpu_memory(func):
    """
    Decorator that ensures proper GPU memory management.

    This wraps functions that might use large amounts of GPU memory,
    making sure memory is properly freed after function execution.
    """
    @functools.wraps(func)

```

```

def wrapper(*args, **kwargs):
    if torch.cuda.is_available():
        # Clear cache before running function
        torch.cuda.empty_cache()
        try:
            return func(*args, **kwargs)
        finally:
            # Clear cache after running function regardless of success/failure
            torch.cuda.empty_cache()
    return func(*args, **kwargs)
return wrapper

#=====
# Step 8: CLIP Model Loading and Evaluation Setup
#=====
# CLIP (Contrastive Language-Image Pre-training) is a neural network that connects
# vision and language. It was trained on 400 million image-text pairs to understand
# the relationship between images and their descriptions.
# We use it here as an "evaluation judge" to assess our generated images.

# Load CLIP model with error handling
try:
    # Load the ViT-B/32 CLIP model (Vision Transformer-based)
    clip_model, clip_preprocess = clip.load("ViT-B/32", device=device)
    print(f"✓ Successfully loaded CLIP model: {clip_model.visual.__class__.__name__}")
except Exception as e:
    print(f"✗ Failed to load CLIP model: {e}")
    clip_available = False
    # Instead of raising an error, we'll continue with degraded functionality
    print("CLIP evaluation will be skipped. Generated images will still be displayed but without quality scores.")

def evaluate_with_clip(images, target_number, max_batch_size=16):
    """
    Use CLIP to evaluate generated images by measuring how well they match textual descriptions.

    This function acts like an "automatic critic" for our generated digits by measuring:
    1. How well they match the description of a handwritten digit
    2. How clear and well-formed they appear to be
    3. Whether they appear blurry or poorly formed

    The evaluation process works by:
    - Converting our images to a format CLIP understands
    - Creating text prompts that describe the qualities we want to measure
    - Computing similarity scores between images and these text descriptions
    - Returning normalized scores (probabilities) for each quality
    """

    Args:
        images (torch.Tensor): Batch of generated images [batch_size, channels, height, width]
        target_number (int): The specific digit (0-9) the images should represent
        max_batch_size (int): Maximum images to process at once (prevents GPU out-of-memory errors)

    Returns:
        torch.Tensor: Similarity scores tensor of shape [batch_size, 3] with scores for:
                    [good handwritten digit, clear digit, blurry digit]
                    Each row sums to 1.0 (as probabilities)
    """
    # If CLIP isn't available, return placeholder scores
    if not clip_available:
        print("⚠️ CLIP not available. Returning default scores.")
        # Equal probabilities (0.33 for each category)
        return torch.ones(len(images), 3).to(device) / 3

    try:
        # For large batches, we process in chunks to avoid memory issues
        # This is crucial when working with big images or many samples
        if len(images) > max_batch_size:
            all_similarities = []

            # Process images in manageable chunks
            for i in range(0, len(images), max_batch_size):
                print(f"Processing CLIP batch {i//max_batch_size + 1}/{(len(images)-1)//max_batch_size + 1}")
                batch = images[i:i+max_batch_size]

                # Use context managers for efficiency and memory management:
                # - torch.no_grad(): disables gradient tracking (not needed for evaluation)

```

```

# - torch.cuda.amp.autocast(): uses mixed precision to reduce memory usage
with torch.no_grad(), torch.cuda.amp.autocast():
    batch_similarities = _process_clip_batch(batch, target_number)
    all_similarities.append(batch_similarities)

# Explicitly free GPU memory between batches
# This helps prevent cumulative memory buildup that could cause crashes
torch.cuda.empty_cache()

# Combine results from all batches into a single tensor
return torch.cat(all_similarities, dim=0)
else:
    # For small batches, process all at once
    with torch.no_grad(), torch.cuda.amp.autocast():
        return _process_clip_batch(images, target_number)

except Exception as e:
    # If anything goes wrong, log the error but don't crash
    print(f"✗ Error in CLIP evaluation: {e}")
    print(f"Traceback: {traceback.format_exc()}")
    # Return default scores so the rest of the notebook can continue
    return torch.ones(len(images), 3).to(device) / 3

def _process_clip_batch(images, target_number):
    """
    Core CLIP processing function that computes similarity between images and text descriptions.

    This function handles the technical details of:
    1. Preparing relevant text prompts for evaluation
    2. Preprocessing images to CLIP's required format
    3. Extracting feature embeddings from both images and text
    4. Computing similarity scores between these embeddings

    The function includes advanced error handling for GPU memory issues,
    automatically reducing batch size if out-of-memory errors occur.

    Args:
        images (torch.Tensor): Batch of images to evaluate
        target_number (int): The digit these images should represent

    Returns:
        torch.Tensor: Normalized similarity scores between images and text descriptions
    """
try:
    # Create text descriptions (prompts) to evaluate our generated digits
    # We check three distinct qualities:
    # 1. If it looks like a handwritten example of the target digit
    # 2. If it appears clear and well-formed
    # 3. If it appears blurry or poorly formed (negative case)
    text_inputs = torch.cat([
        clip.tokenize(f"A handwritten number {target_number}"),
        clip.tokenize(f"A clear, well-written digit {target_number}"),
        clip.tokenize(f"A blurry or unclear number")
    ]).to(device)

    # Process images for CLIP, which requires specific formatting:

    # 1. Handle different channel configurations (dataset-dependent)
    if IMG_CH == 1:
        # CLIP expects RGB images, so we repeat the grayscale channel 3 times
        # For example, MNIST/Fashion-MNIST are grayscale (1-channel)
        images_rgb = images.repeat(1, 3, 1, 1)
    else:
        # For RGB datasets like CIFAR-10/CelebA, we can use as-is
        images_rgb = images

    # 2. Normalize pixel values to [0,1] range if needed
    # Different datasets may have different normalization ranges
    if images_rgb.min() < 0: # If normalized to [-1,1] range
        images_rgb = (images_rgb + 1) / 2 # Convert to [0,1] range

    # 3. Resize images to CLIP's expected input size (224x224 pixels)
    # CLIP was trained on this specific resolution
    resized_images = F.interpolate(images_rgb, size=(224, 224),
                                   mode='bilinear', align_corners=False)

```

```

# Extract feature embeddings from both images and text prompts
# These are high-dimensional vectors representing the content
image_features = clip_model.encode_image(resized_images)
text_features = clip_model.encode_text(text_inputs)

# Normalize feature vectors to unit length (for cosine similarity)
# This ensures we're measuring direction, not magnitude
image_features = image_features / image_features.norm(dim=-1, keepdim=True)
text_features = text_features / text_features.norm(dim=-1, keepdim=True)

# Calculate similarity scores between image and text features
# The matrix multiplication computes all pairwise dot products at once
# Multiplying by 100 scales to percentage-like values before applying softmax
similarity = (100.0 * image_features @ text_features.T).softmax(dim=-1)

return similarity

```

```

except RuntimeError as e:
    # Special handling for CUDA out-of-memory errors
    if "out of memory" in str(e):
        # Free GPU memory immediately
        torch.cuda.empty_cache()

        # If we're already at batch size 1, we can't reduce further
        if len(images) <= 1:
            print("✖ Out of memory even with batch size 1. Cannot process.")
            return torch.ones(len(images), 3).to(device) / 3

        # Adaptive batch size reduction - recursively try with smaller batches
        # This is an advanced technique to handle limited GPU memory gracefully
        half_size = len(images) // 2
        print(f"⚠️ Out of memory. Reducing batch size to {half_size}.")

        # Process each half separately and combine results
        # This recursive approach will keep splitting until processing succeeds
        first_half = _process_clip_batch(images[:half_size], target_number)
        second_half = _process_clip_batch(images[half_size:], target_number)

        # Combine results from both halves
        return torch.cat([first_half, second_half], dim=0)

    # For other errors, propagate upward
    raise e

```

```

=====
# CLIP Evaluation - Generate and Analyze Sample Digits
=====

# This section demonstrates how to use CLIP to evaluate generated digits
# We'll generate examples of all ten digits and visualize the quality scores

try:
    for number in range(10):
        print(f"\nGenerating and evaluating number {number}...")

        # Generate 4 different variations of the current digit
        samples = generate_number(model, number, n_samples=4)

        # Evaluate quality with CLIP (without tracking gradients for efficiency)
        with torch.no_grad():
            similarities = evaluate_with_clip(samples, number)

        # Create a figure to display results
        plt.figure(figsize=(15, 3))

        # Show each sample with its CLIP quality scores
        for i in range(4):
            plt.subplot(1, 4, i+1)

            # Display the image with appropriate formatting based on dataset type
            if IMG_CH == 1: # Grayscale images (MNIST, Fashion-MNIST)
                plt.imshow(samples[i][0].cpu(), cmap='gray')
            else: # Color images (CIFAR-10, CelebA)
                img = samples[i].permute(1, 2, 0).cpu() # Change format for matplotlib
                if img.min() < 0: # Handle [-1,1] normalization
                    img = (img + 1) / 2 # Convert to [0,1] range

```

```

# Convert to tensor
plt.imshow(img)

# Extract individual quality scores for display
# These represent how confidently CLIP associates the image with each description
good_score = similarities[i][0].item() * 100 # Handwritten quality
clear_score = similarities[i][1].item() * 100 # Clarity quality
blur_score = similarities[i][2].item() * 100 # Blurriness assessment

# Color-code the title based on highest score category:
# - Green: if either "good handwritten" or "clear" score is highest
# - Red: if "blurry" score is highest (poor quality)
max_score_idx = torch.argmax(similarities[i]).item()
title_color = 'green' if max_score_idx < 2 else 'red'

# Show scores in the plot title
plt.title(f'Number {number}\nGood: {good_score:.0f}\nClear: {clear_score:.0f}\nBlurry: {blur_score:.0f}', color=title_color)
plt.axis('off')

plt.tight_layout()
plt.show()
plt.close() # Properly close figure to prevent memory leaks

# Clean up GPU memory after processing each number
# This is especially important for resource-constrained environments
torch.cuda.empty_cache()

except Exception as e:
    # Comprehensive error handling to help students debug issues
    print(f"X Error in generation and evaluation loop: {e}")
    print("Detailed error information:")
    import traceback
    traceback.print_exc()

    # Clean up resources even if we encounter an error
    if torch.cuda.is_available():
        print("Clearing GPU cache...")
        torch.cuda.empty_cache()

#=====
# STUDENT ACTIVITY: Exploring CLIP Evaluation
#=====
# This section provides code templates for students to experiment with
# evaluating larger batches of generated digits using CLIP.

print("\nSTUDENT ACTIVITY:")
print("Try the code below to evaluate a larger sample of a specific digit")
print("""
# Example: Generate and evaluate 10 examples of the digit 6
# digit = 6
# samples = generate_number(model, digit, n_samples=10)
# similarities = evaluate_with_clip(samples, digit)
#
# # Calculate what percentage of samples CLIP considers "good quality"
# # (either "good handwritten" or "clear" score exceeds "blurry" score)
# good_or_clear = (similarities[:,0] + similarities[:,1] > similarities[:,2]).float().mean()
# print(f"CLIP recognized {good_or_clear.item()}*100.1f% of the digits as good examples of {digit}")
#
# # Display a grid of samples with their quality scores
# plt.figure(figsize=(15, 8))
# for i in range(len(samples)):
#     plt.subplot(2, 5, i+1)
#     plt.imshow(samples[i][0].cpu(), cmap='gray')
#     quality = "Good" if similarities[i,0] + similarities[i,1] > similarities[i,2] else "Poor"
#     plt.title(f"Sample {i+1}: {quality}", color='green' if quality == "Good" else 'red')
#     plt.axis('off')
# plt.tight_layout()
# plt.show()
""")

```

